

PRODUCTIVE YOU

LEAD GENERATION BUSINESS OPERATING SYSTEM

The Complete Intelligence Architecture for Research, Lead Generation,
Goal Execution & Agency Operations

For: Software Dev Agency Pivoting to AI Automation

Version: 2.0 — Full System Blueprint

Date: February 2026

□ TABLE OF CONTENTS

1. PART 1: GOAL OPERATING SYSTEM	p. 3
2. PART 2: RESEARCH ENGINE — INTELLIGENT & CACHED	p. 6
3. PART 3: LEAD GENERATION PIPELINE	p. 9
4. PART 4: AUTOMATED LEAD QUALIFICATION SYSTEM	p. 12
5. PART 5: AGENCY PACKAGES & PRODUCTIZED SERVICES	p. 15
6. PART 6: CRM & SALES WORKFLOW SOP	p. 17
7. PART 7: 90-DAY EXECUTION PLAN	p. 20
8. PART 8: MASTER PROMPTS & AI TEMPLATES	p. 23

1.1 The Goal Architecture

Every goal in Productive You lives inside a structured hierarchy that cascades from your biggest vision down to what you do today. This is not a static list — it is a living, AI-assisted execution engine.

LAYER	MODEL	WHAT LIVES HERE	AI ROLE
YEAR	YearPlan	Vision, theme, focus areas, 3–5 goals	AI drafts your year vision based on business context
QUARTER	QuarterPlan + QuarterFocus	3 Big Objectives + Goal Focus areas	AI suggests quarterly priorities based on goal gap
MONTH	MonthPlan + MonthFocus	Monthly objectives, theme, milestones	AI reviews progress and recommends next actions
WEEK	WeekPlan + WeekFocus	Top outcomes, planned hours, key wins	AI generates optimal weekly schedule from tasks
DAY	DayPlan + TimeBlocks	Time blocks, priorities, energy tracking	AI allocates tasks by energy level + deadline

1.2 Milestone & Pivot Tracking (Your Make.com Example)

When you are pivoting — like learning Make.com automation — every new skill or business capability is a Milestone inside a Goal. Here is exactly how this maps:

ELEMENT	YOUR PIVOT EXAMPLE
Goal	"Pivot software agency to AI Automation & Lead Gen agency by Q4 2026"
Milestone 1	Master Make.com — build 3 live automation workflows
Milestone 2	Launch first Lead Gen productized service (\$500/mo)
Milestone 3	Sign 5 retainer clients for AI-powered lead gen
KeyStep (under M1)	Complete Make.com fundamentals course → Build Zap → Build Lead enrichment scenario
Task	Daily: "Watch Make.com Module 3" — 45 min time block — Tuesday 9AM

1.3 AI Research Per Milestone

□ HOW IT WORKS

When you click into any Milestone (e.g., "Master Make.com"), you can trigger the Research Engine with a milestone-specific prompt. The AI scrapes the web, reads forums, documentation, and expert breakdowns, then returns 10 Actionable Insights and auto-generates Task recommendations directly linked to that milestone.

Example Milestone Research Prompt: *"I am a software developer pivoting to AI automation. What are the most critical things I need to learn in Make.com to build lead generation and CRM automation workflows for B2B clients in 2026? Give me a prioritized learning path and specific scenarios to build."*

1.4 The Intelligent Calendar — Google-Grade

The calendar must be the command center of daily execution. Every task, milestone, time block, and meeting appears here. The design requirement:

VIEW	WHAT IT SHOWS	INTERACTION
YEAR	Milestones, goal deadlines, quarter themes as banners	Click month → drill into MonthPlan
MONTH	Weekly themes, key tasks, client meetings, energy blocks	Click week → WeekPlan view with drag-drop scheduling
WEEK	Time blocks by hour (like Google Cal), color-coded by type	Drag to reschedule, click to edit, AI to optimize
DAY	Hour-by-hour execution view with energy meter, timer	One-click timer start, add ad-hoc tasks, AI reschedule

Mobile-First Calendar Requirements: Swipe left/right to navigate days. Bottom tab for Day/Week/Month/Year. Floating "+" button to add task or event. Pull-to-refresh for AI schedule recalculation. Color-coded task chips (red = urgent, blue = client, green = personal, gold = milestone).

2.1 Current Problems & Fixes

PROBLEM	IMPACT	SOLUTION
Re-scraping same URLs on repeated queries	Token waste, slow results, credit burn	URL-level cache with 7-day TTL in Redis/DB
No deduplication of insights	LLM re-processes identical content	Content hash fingerprinting before LLM call
No query similarity check	Slight rephrasing = full new research	Semantic embedding similarity — reuse if >85% match
Full page content sent to LLM	Expensive, context overflow	Jina extracts → chunk → only relevant sections sent
No source credibility filter	Low-quality sources pollute insights	Domain authority score filter before scraping

2.2 Caching Architecture — Industry Standard

CACHE LAYER	WHAT IT STORES	TTL	IMPLEMENTATION
L1: Query	Full research result for identical prompts	24 hours	Redis key: SHA256(userId + normalizedPrompt)
L2: URL	Scraped markdown content per URL	7 days	DB: ResearchSource.content + scrapedAt check
L3: Semantic	Embeddings of past research queries	30 days	Pgvector: cosine similarity >0.85 = cache hit
L4: Domain	Domain authority + credibility scores	90 days	Moz/Ahrefs API or local DA score table

2.3 The Intelligent Research Pipeline — Full Flow

#	STEP	DETAIL
1	Query Normalization	Strip stopwords, lowercase, extract intent keywords. Generate SHA256 hash for L1 cache lookup.
2	Cache Check (L1)	If identical query found in Redis within TTL → return cached result immediately. ZERO tokens used.

3	Semantic Similarity (L3)	Embed query with text-embedding-3-small. Query pgvector for similar past researches. If >85% match, return enriched cached result.
4	Search Query Generation	LLM generates 5–8 targeted Google search queries optimized for B2B lead intelligence (see query templates below).
5	Serper API Search	Execute searches. Collect top 10 URLs per query. Filter: remove domains with DA <20, social media, job boards.
6	URL Dedup + Cache Check (L2)	Check each URL against ResearchSource table. Skip if scraped within 7 days — reuse stored content.
7	Jina Scraping (New URLs Only)	Scrape only un-cached URLs. Extract clean markdown. Chunk into 800-token segments. Store in ResearchSource.
8	Relevance Scoring	Score each chunk 0–1 for relevance to research goal. Only chunks >0.6 sent to LLM. Massive token savings.
9	LLM Insight Extraction	Feed scored chunks to Gemini 2.0 Flash. Extract structured ResearchInsight records: title, content, category, confidence, visualData.
10	Action Item Generation	Second LLM pass: "Given these insights, generate 10 concrete next actions for a B2B agency." Save as ActionItem records.
11	Lead Extraction (if LEAD_GENERATION scope)	Third LLM pass: extract structured leads with Name, Company, Email guess, Industry, Pain Points, suggested DM + Email.
12	Cache Write	Write full result to L1 Redis cache. Embeddings saved for L3.

2.4 Guaranteed Lead-Rich Search Query Templates

These search query patterns are battle-tested for B2B lead intelligence. The research engine should auto-generate variations of these:

INTENT	SEARCH QUERY TEMPLATE
Company Discovery	"[industry] companies in [location] 2025 site:linkedin.com OR site:clutch.co"
Pain Point Mining	"[industry] challenges 2025 OR problems OR struggling with automation"
Decision Maker Find	"[industry] CEO OR CTO OR "Head of" OR "VP of" site:linkedin.com [location]"
Tech Stack Intel	"companies using [tech] [industry] [location] -jobs -careers"
Budget Signal	"[industry] [location] hiring automation OR digital transformation 2025"
Competitor Clients	"[competitor name] clients OR case studies OR success stories [industry]"
Intent Signals	"[industry] looking for agency OR outsource OR partner [location] 2025"
Directory Mining	"[industry] directory [location] site:clutch.co OR site:g2.com OR site:capterra.com"

3.1 Lead Data Architecture

Every lead generated flows through a qualification funnel before it ever reaches a human. Here is the complete data model and what each field should contain:

FIELD	SOURCE	PURPOSE IN QUALIFICATION
name	LLM extraction	Primary contact — drives personalization in outreach
company	LLM + Serper	Company-level research triggers enrichment pipeline
email	LLM guess / Hunter.io / Apollo	Verified before outreach — invalid = auto-disqualified
phone	LLM / Clearbit	High-intent channel — reserved for qualified leads only
website	Serper result	Feeds tech stack detection (BuiltWith API) for personalization
industry	LLM + SIC code	Industry-level scoring — some industries worth 3x more
location	LLM extraction	Time zone routing + local outreach compliance
painPoints	LLM from content	Used to select email template + DM angle — most critical field
suggestedDM	LLM generated	AI-written LinkedIn message — personalized to pain points
suggestedEmail	LLM generated	AI-written cold email — ready to send after enrichment
personalization	LLM JSON	Company news, recent posts, tech signals for hyper-personalization
contacted	System tracked	Prevents duplicate outreach — gate before CrmLead creation

3.2 The Lead Generation Workflow — Step by Step

PHASE	ACTION	SYSTEM BEHAVIOUR
TRIGGER	User sets ResearchScope = LEAD_GENERATION	Ingest event fired with lead-specific pipeline flag
SEARCH	Research Engine runs 8 targeted queries	Serper returns top B2B prospect URLs, directories, LinkedIn
SCRAPE	Jina extracts company pages	Pulls About pages, team pages, LinkedIn company profiles

EXTRACT	Gemini extracts lead JSON array	Structured: name, company, industry, pain points, signals
ENRICH	Hunter.io / Apollo API called	Verifies email, adds phone, LinkedIn URL, company size, revenue
SCORE	Lead Scoring Engine runs	Score 0-100 based on fit matrix (see 3.3)
QUALIFY	Threshold filter applied	Score $\geq 60 \rightarrow$ auto-moved to CrmLead as "QUALIFIED". $<60 \rightarrow$ stays in Lead table for review
PERSONALIZE	GPT-4o mini second pass	Rewrites suggestedEmail and suggestedDM with enriched data
REVIEW	User reviews in Lead Dashboard	One-click to push to CRM Pipeline \rightarrow Discovery stage

4.1 Lead Scoring Matrix — The Fit Score

Every lead gets an automated score from 0–100. This score determines if the lead enters your CRM, gets emailed automatically, or is discarded. Customize thresholds per campaign.

SCORING DIMENSION	MAX PTS	WEIGHT	HOW SCORED
Industry Fit	25	25%	Your target industries = 25pts. Adjacent = 15. Irrelevant = 0.
Company Size Fit	20	20%	Ideal ICP size (e.g., 10–200 employees) = 20pts. Too small/large = 5pts.
Pain Point Match	20	20%	LLM checks if extracted pain points match your service. NLP similarity score × 20.
Email Verified	15	15%	Hunter.io confidence ≥80% = 15pts. Guessed = 5pts. No email = 0.
Tech Stack Signal	10	10%	Uses tools you automate (e.g., HubSpot, Salesforce) = 10pts.
Buying Intent Signal	10	10%	Recent job postings, growth news, or hiring in target area = 10pts.

SCORE RANGE	TIER	STATUS	AUTOMATED ACTION
80–100	□ HOT	QUALIFIED	Auto-create CrmLead → Discovery stage + trigger email sequence within 24h
60–79	✓WARM	REVIEW	Create CrmLead → Notify user to review → Hold for manual send approval
40–59	□ COLD	NURTURE	Save to Lead table only → Add to long-term drip list → Re-score in 30 days
0–39	X DISCARD	UNQUALIFIED	Mark as UNQUALIFIED → Archive → Log reason for model improvement

4.2 CRM Workflow Automations — The Inngest Functions

These are the Inngest workflow triggers that run automatically in your system:

TRIGGER	CONDITION	ACTION SEQUENCE
---------	-----------	-----------------

LEAD_CREATED	Score ≥ 80	Create CrmLead → Assign to Discovery → Schedule email Day 1, 4, 8 → Create follow-up Task
LEAD_SCORE_THRESHOLD	Score crosses 60 after enrichment	Promote from Lead table → CrmLead → Notify user with context card
INACTIVITY_ALERT	No engagement in 7 days	Move to Nurture tag → Schedule re-engagement email → Create task reminder
LEAD_STATUS_CHANGED	Status = QUALIFIED	Trigger proposal template email → Create Meeting booking task → Update pipeline stage
DEAL_STAGE_CHANGED	Moves to Proposal	Send proposal email template → Create DealNote → Schedule follow-up in 3 days
DEAL_CLOSED_WON	Status = WON	Create Client record → Create onboarding project + tasks → Archive CrmLead → Log win
SCHEDULED	Daily 6AM cron	Re-score stale leads → Check inactivity → Send digest of hot leads to user

05

AGENCY PACKAGES & PRODUCTIZED SERVICES

Three clear offers you can sell immediately

<input type="checkbox"/> STARTER — THE COLD EMAIL ENGINE	<input type="checkbox"/> GROWTH — AI SCRAPING LEADS	<input type="checkbox"/> SCALE — FULL AUTOMATION SYSTEM
\$500–\$800 / month Deliverables: <ul style="list-style-type: none">• 50 verified leads/month• 3 email sequences written by AI• CRM pipeline setup• Weekly lead report Research Engine Use: <ul style="list-style-type: none">• 1 research run per week• LEAD_GENERATION scope• Serper + Jina pipeline	\$1,500–\$2,500 / month Deliverables: <ul style="list-style-type: none">• 200 enriched leads/month• AI-personalized email + DM per lead• Full CRM with scoring + routing• Competitor intelligence report• Bi-weekly strategy call Research Engine Use: <ul style="list-style-type: none">• 3 research runs per week• Apollo.io enrichment layer• Automated email sequencing	\$4,000–\$8,000 / month Deliverables: <ul style="list-style-type: none">• • • • • • • Research Engine Use: <ul style="list-style-type: none">• • •

6.1 The Complete Client Acquisition Flow

DAY	STAGE	ACTIONS IN PRODUCTIVE YOU	AUTOMATION
D0	Lead Created	Research Engine outputs lead. Scoring engine qualifies. CrmLead created in Discovery stage.	Auto: Ingest LEAD_CREATED event fires. Email Day 1 queued.
D1	First Touch	Email 1 sent (AI-personalized). LinkedIn DM sent. CrmActivity logged automatically.	Auto: SendGrid sends cold email. CrmLeadCommunication record created.
D4	Follow-Up 1	Email 2 sent (value-add, industry insight). Review open/click data in CrmLead record.	Auto: If no open → use different subject line variant. If opened → accelerate sequence.
D8	Follow-Up 2	Email 3 (case study or social proof). If no reply → move to Nurture tag.	Auto: Tag added. Moved to 30-day drip sequence.
D12	Reply Received	Lead status → CONTACTED. Create Meeting task. Book discovery call link sent.	Manual: User responds. System logs sentiment score on communication.
D14	Discovery Call	Meeting logged with agenda. AI summarizes call notes. Action items auto-created.	Post-call: AI generates proposal outline based on call notes.
D17	Qualification	Update CrmLead with pain points, budget, timeline. Score recalculated. Move to QUALIFIED.	Auto: If score ≥80 after call → Proposal email template queued.
D20	Proposal	Deal created in Pipeline. Proposal PDF sent. DealNote logged. Follow-up task created.	Auto: Proposal stage email triggers. 3-day follow-up reminder set.
D28	Negotiation	Counter-offers logged as DealNote. Pipeline stage updated. Probability adjusted.	Manual review. AI suggests pricing flexibility based on deal value.
D35	CLOSED WON	Client record created from CrmLead. Onboarding project auto-created with tasks.	Auto: DEAL_CLOSED_WON triggers full onboarding workflow.

07**90-DAY EXECUTION PLAN***From pivot to first 5 agency clients signed***\$0→\$5K**

MRR Target

5

Clients to Sign

90

Days to Execute

3

Quarter Objectives

7.1 Quarter Objectives (QuarterFocus)

OBJ	TITLE	SUCCESS METRIC
QO-1	Master AI Automation Stack	3 live Make.com workflows built + certified in Zapier + Make
QO-2	Launch Productized Lead Gen Service	Service page live + 10 cold outreach campaigns running
QO-3	Sign 5 Paying Agency Clients	\$3,000+ MRR from lead gen retainers

7.2 Month-by-Month Breakdown**MONTH 1 — FOUNDATION (Days 1–30)**

WEEK	FOCUS	DELIVERABLE
W1	Make.com fundamentals + first scenario	1 live automation: Lead captured → CRM → Email sent
W2	Optimize Research Engine caching	L1 + L2 cache implemented. Token cost reduced 60%.
W3	Build lead scoring engine	Automated scoring matrix live. 50 test leads scored.
W4	Launch first cold outreach campaign	100 leads generated. First campaign live. Track opens.

MONTH 2 — TRACTION (Days 31–60)

WEEK	FOCUS	DELIVERABLE
W5	Apollo.io / Hunter.io enrichment integration	Email verification live. Bounce rate <5%.
W6	First 2 discovery calls booked	Sales SOP tested. Proposal template created.
W7	LinkedIn outreach via PhantomBuster	50 LinkedIn DMs sent/week. Connection acceptance tracked.

W8	Close first 2 clients	\$1,000–\$2,000 MRR. Onboarding workflow activated.
----	-----------------------	---

MONTH 3 — SCALE (Days 61–90)

WEEK	FOCUS	DELIVERABLE
W9	Client portal MVP launch	Clients can view their pipeline in real-time.
W10	Systemize delivery — SOPs documented	Repeatable delivery. Can serve 10+ clients without new tools.
W11	Close clients 3, 4, 5	Referral strategy activated. Testimonials collected.
W12	\$5K+ MRR celebration + Q2 planning	Quarter review. Set next quarter goal: \$15K MRR.

8.1 Goal Research Prompt

PROMPT: New Skill / Pivot Milestone Research

Act as an expert business strategist and learning architect. I am [DESCRIBE YOUR BACKGROUND]. I am pivoting to [PIVOT DESCRIPTION]. My specific goal is: [GOAL]. For the milestone "[MILESTONE NAME]", I need: (1) A prioritized, step-by-step learning path with specific resources, (2) The 5 most common mistakes people make when doing this, (3) The fastest path to a first paying client using this skill, (4) The exact tools and tech stack I need, (5) 10 concrete action items I can start this week. Format as structured insights I can act on immediately.

8.2 Lead Generation Research Prompt

PROMPT: B2B Lead Generation Research

Act as a B2B lead generation specialist. I need you to identify and structure potential business leads for the following: TARGET INDUSTRY: [industry]. TARGET LOCATION: [location/remote]. COMPANY SIZE: [e.g., 10-200 employees]. MY SERVICE: [your service offering]. PAIN POINTS I SOLVE: [list]. For each lead you find, provide: Company name, Decision maker name and title, Estimated email pattern, Industry and company size, Top 3 pain points based on their online presence, A personalized cold email subject line (under 50 chars), A personalized LinkedIn DM opener (under 300 chars), Confidence score 0-100 that they need my service. Return as a structured JSON array.

8.3 Lead Qualification Prompt

PROMPT: Lead Scoring & Qualification

Act as a sales qualification expert. I have a lead with the following data: [INSERT LEAD JSON]. My ideal client profile is: Industry: [industry], Size: [size], Budget: [budget range], Pain points: [list]. Score this lead 0-100 on these dimensions: (1) Industry fit 0-25, (2) Company size fit 0-20, (3) Pain point alignment 0-20, (4) Contact quality 0-15, (5) Tech stack signals 0-10, (6) Buying intent signals 0-10. Return: total score, tier (HOT/WARM/COLD/DISCARD), top 3 personalization hooks, recommended outreach channel, suggested next action, and reasoning for the score.

8.4 Competitive Intelligence Prompt

PROMPT: Competitor & Market Intelligence

Act as a market research analyst specializing in B2B SaaS and agency markets. Research the competitive landscape for [YOUR SERVICE] in [MARKET/LOCATION]. I need: (1) Top 5 direct competitors with their pricing, positioning, and weaknesses, (2) Top 5 indirect competitors I am fighting for attention against, (3) Market size estimate and growth trajectory for 2025-2026, (4) The #1 underserved pain point in this market that competitors are missing, (5) The exact positioning statement I should use to differentiate, (6) Top 3 acquisition channels that are working in this space right now, (7) 5 companies currently winning in this space and WHY. Make this actionable for someone building a new agency from scratch.

8.5 Daily Execution Prompt (AI Day Planner)

PROMPT: Intelligent Daily Schedule

Act as my AI Chief of Staff and productivity optimizer. Today is [DATE]. My energy levels are: Morning [1-10], Afternoon [1-10], Evening [1-10]. My top priorities this week are: [LIST]. My must-do tasks today are: [LIST WITH ESTIMATES]. My optional tasks: [LIST]. My meetings: [LIST WITH TIMES]. Design my optimal day plan with: (1) Hour-by-hour time blocks (8AM-6PM), (2) Tasks allocated to energy levels (deep work in peak energy, admin in low energy), (3) Buffer time for unexpected work, (4) A concrete evening wind-down and tomorrow preparation. Format as a clean schedule I can follow immediately.

PRODUCTIVE YOU — LEAD GENERATION BUSINESS OPERATING SYSTEM v2.0

Built for Agency Founders Who Build Systems, Not Just Services.